

Design Space of Graph Neural Networks

CS224W: Machine Learning with Graphs

Charilaos Kanatsoulis, Stanford University

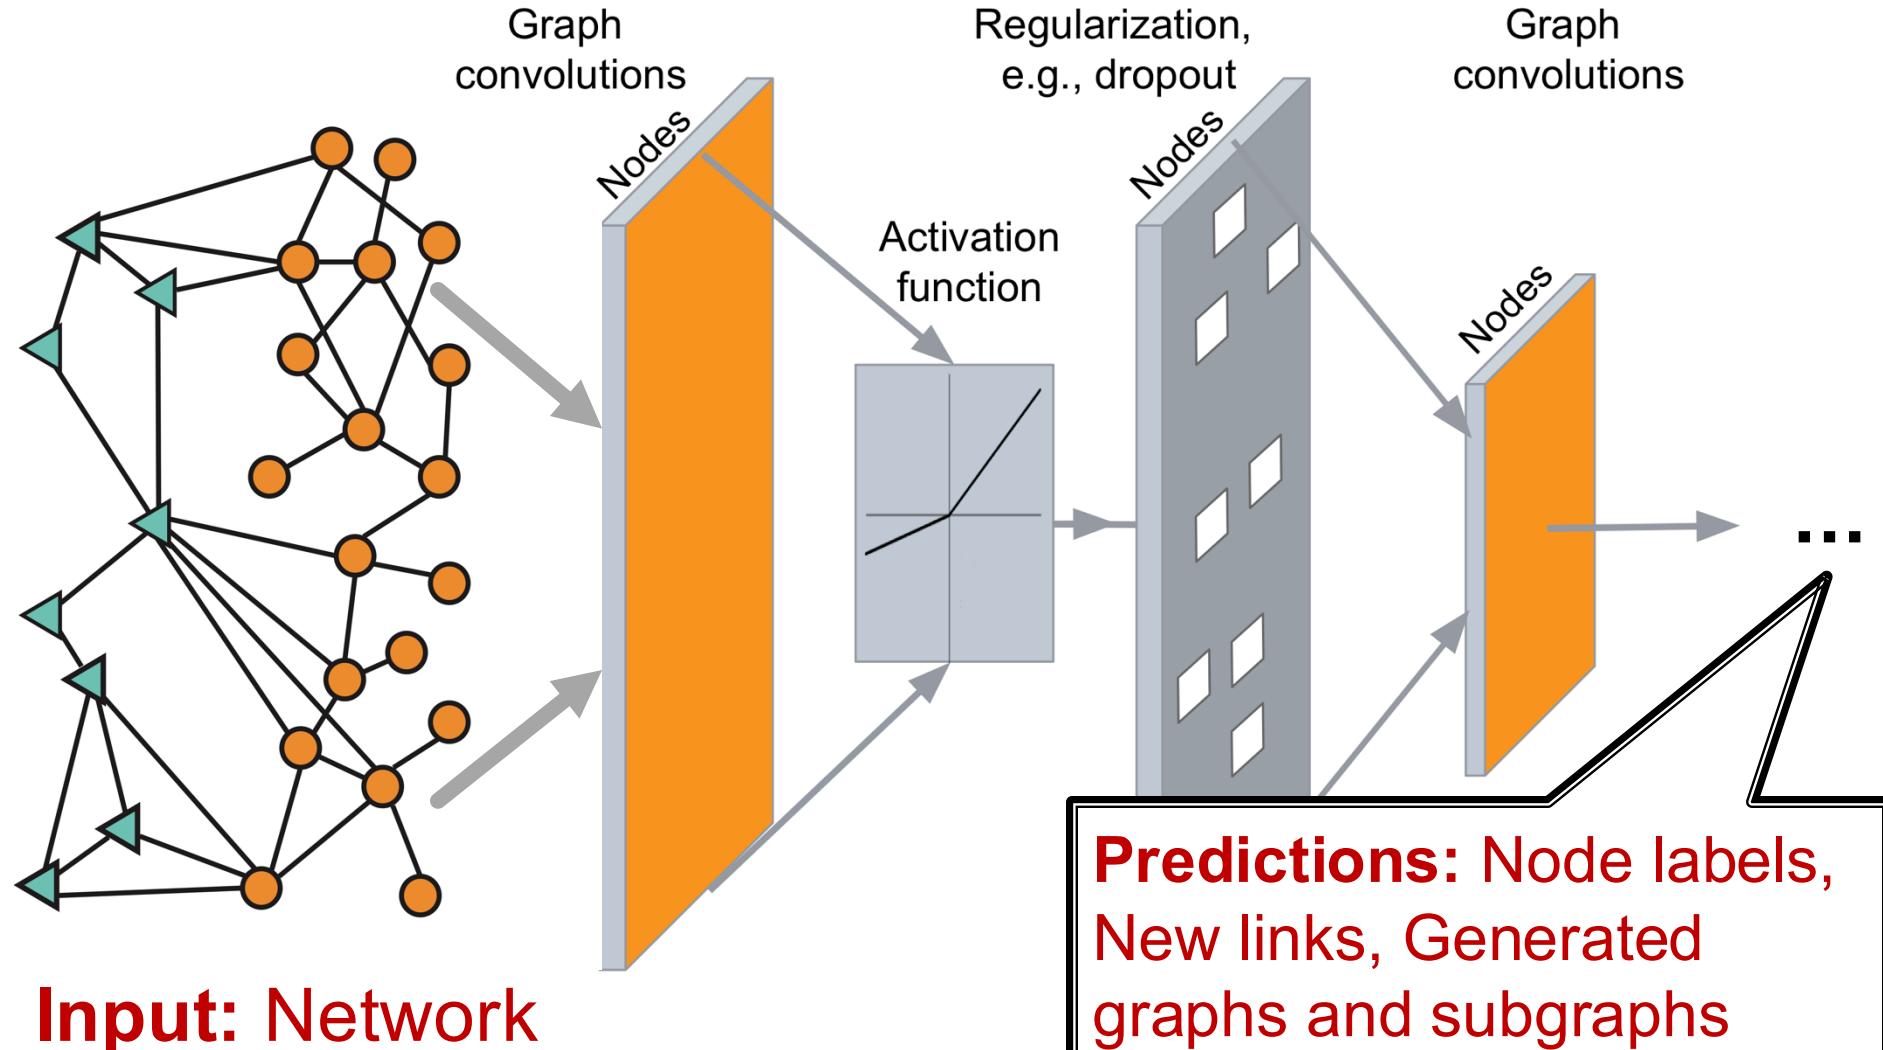
<http://cs224w.stanford.edu>



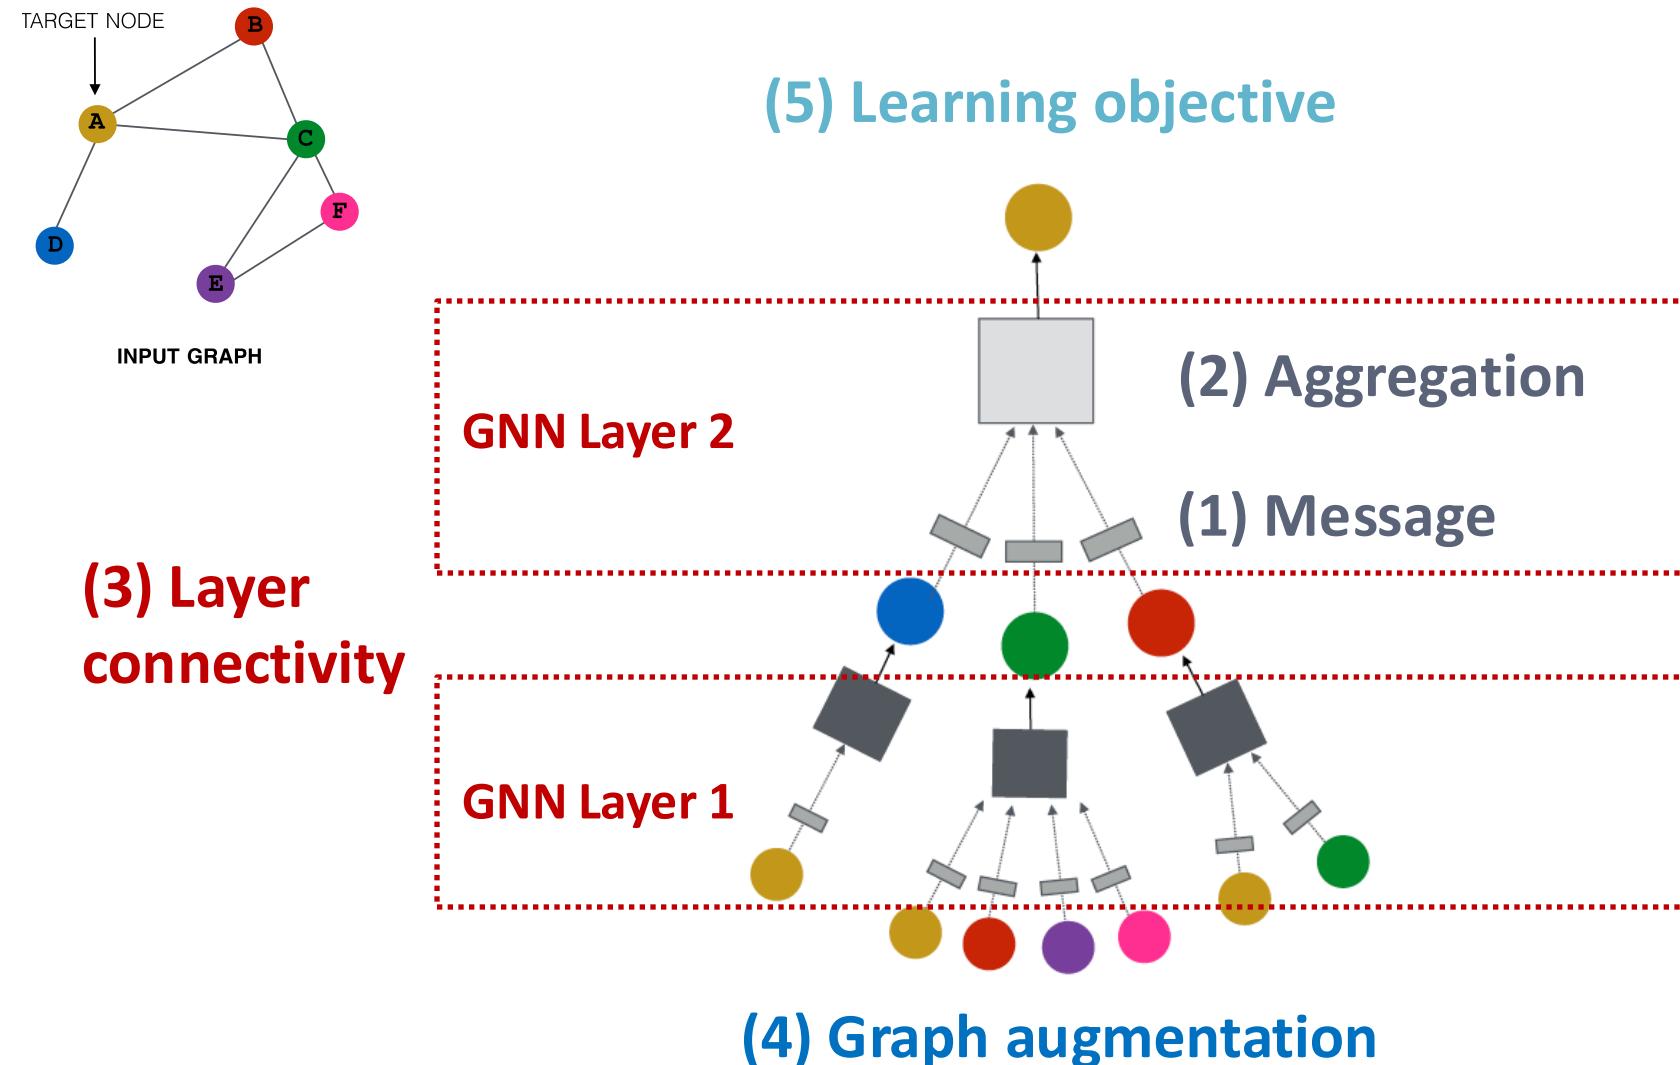
Announcements

- Congratulations on completing this class! We are happy that you joined us.
 - Today is the last lecture!
- **Exam grades have been released!**
 - Regrade requests open until 12/10 EOD
- **Colab 5** due today (12/4)
- **Project Report** due next Thursday (12/11)
 - We will open a form on Gradescope to upload a link to your Medium article and code

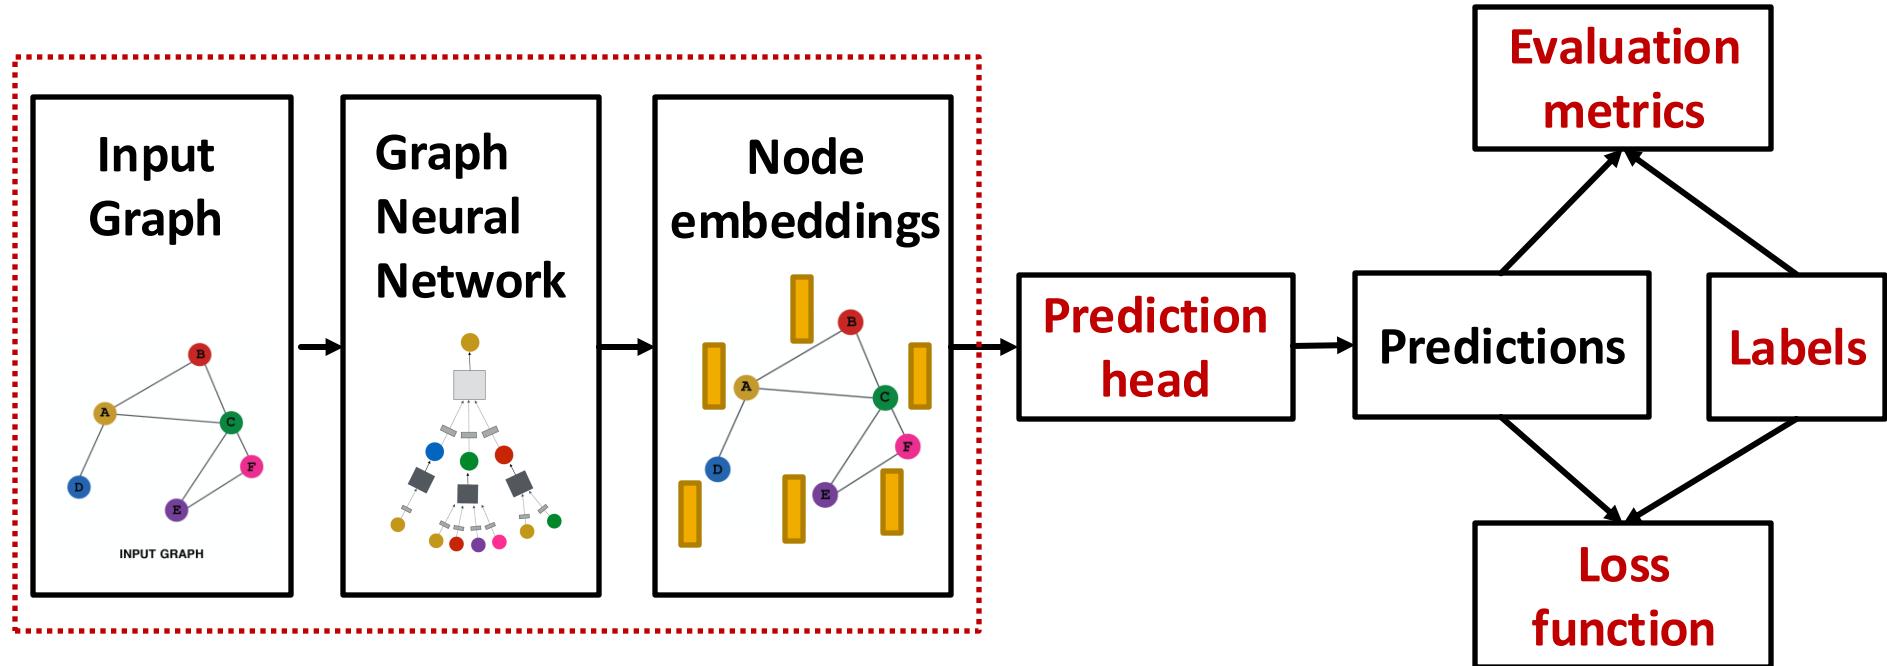
CS224W: Deep Learning in Graphs



A General GNN Framework



GNN Training Pipeline



Output of a GNN: set of node embeddings

$$\{\mathbf{h}_v^{(L)}, \forall v \in G\}$$

Key Questions for GNN Design

- **GNN architectural design:**
 - How to find a good GNN design for a specific GNN task?
- **Important but challenging:**
 - Domain experts want to use SOTA GNN on their specific tasks, however...
 - There are tons of possible GNN architectures
 - GCN, GraphSAGE, GAT, GIN, ...
 - **Issue:** Best design in one task can perform badly for another task
 - Redo hyperparameter grid search for each new task is NOT feasible
- **Topic for today:**
 - Study for the ***GNN design space and task space***
 - **GraphGym**, a powerful platform for exploring different GNN designs and tasks

Background: Terminology

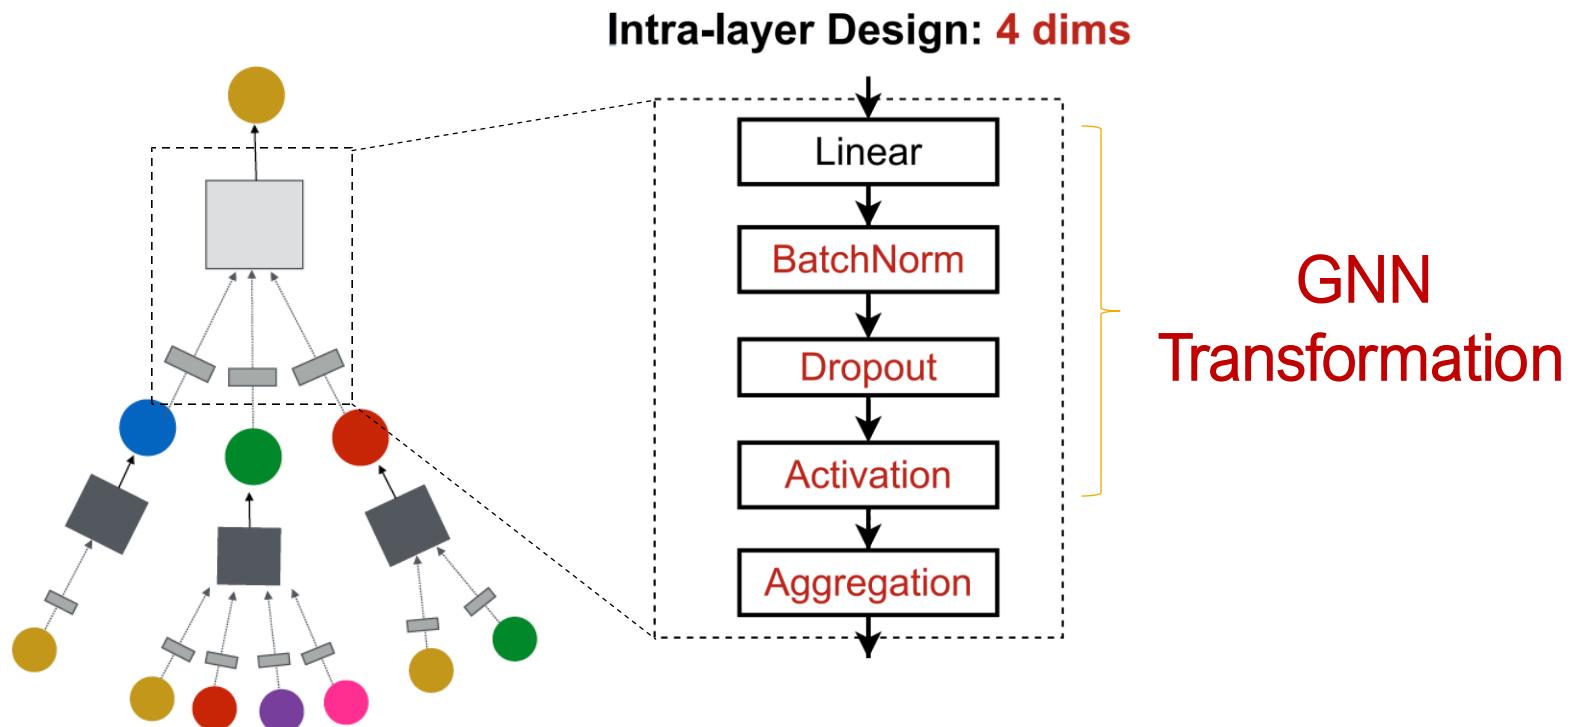
- **Design:** a concrete model instantiation
 - E.g., a 4-layer GraphSAGE
- **Design dimensions** characterize a design
 - E.g., the number of layers $L \in \{2, 4, 6, 8\}$
- **Design choice** is the actual selected value in the design dimension
 - E.g., the number of layers $L = 2$
- **Design space** consists of a Cartesian product of design dimensions
- **Task:** A specific task of interest
 - E.g., node classification on Cora, graph classification on ENZYMEs
- **Task space** consists of all the tasks we care about

Recap: GNN Design Space

Intra-layer Design:

GNN Layer = Transformation + Aggregation

- We propose a general instantiation under this perspective

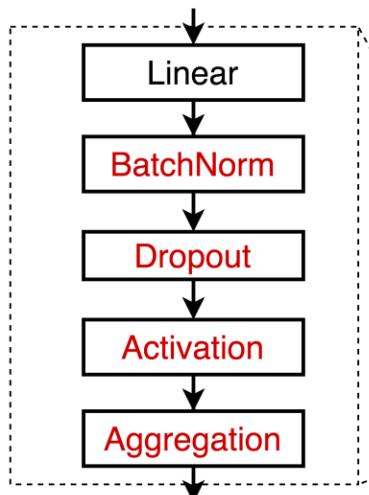


Recap: GNN Design Space

Inter-layer Design

- We explore different ways of organizing GNN layers

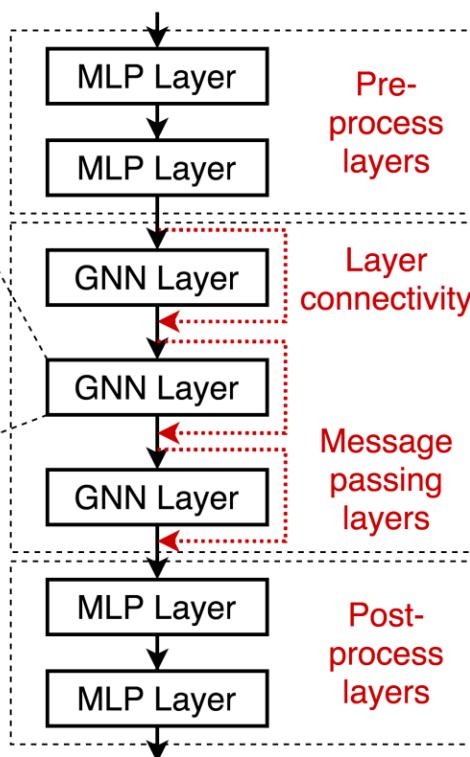
Intra-layer Design: 4 dims



Learning Configuration: 4 dims

Batch size
Learning rate
Optimizer
Training epochs

Inter-layer Design: 4 dims



Pre-process layers:

Important when expressive node feature encoder is needed
E.g., when nodes are images/text

Skip connections:

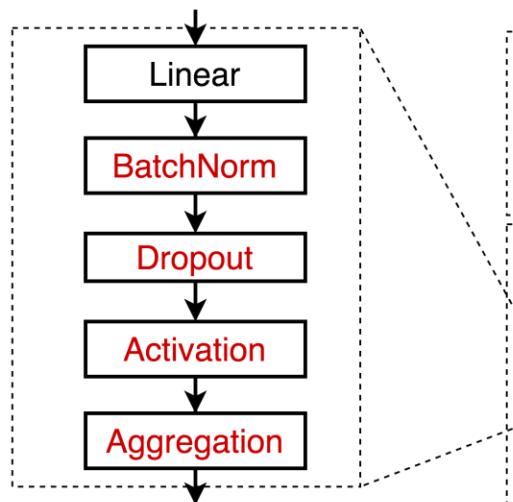
Improve deep GNN's performance

Post-process layers:

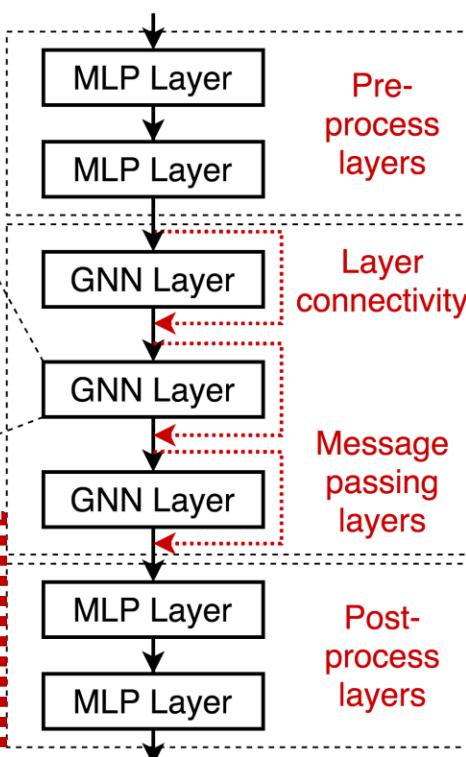
Important when reasoning or transformation over node embeddings are needed
E.g., graph classification, knowledge graphs

Recap: GNN Design Space

Intra-layer Design: 4 dims



Inter-layer Design: 4 dims



Learning Configuration: 4 dims

Batch size
Learning rate
Optimizer
Training epochs

Learning configurations

- Often neglected in current literature
- But we found they have high impact on performance

Summary: GNN Design Space

■ Overall: A GNN design space

■ Intra-layer design

Batch Normalization	Dropout	Activation	Aggregation
True, False	False, 0.3, 0.6	RELU, PRELU, SWISH	MEAN, MAX, SUM

■ Inter-layer design

Layer connectivity	Pre-process layers	Message passing layers	Post-precess layers
STACK, SKIP-SUM, SKIP-CAT	1, 2, 3	2, 4, 6, 8	1, 2, 3

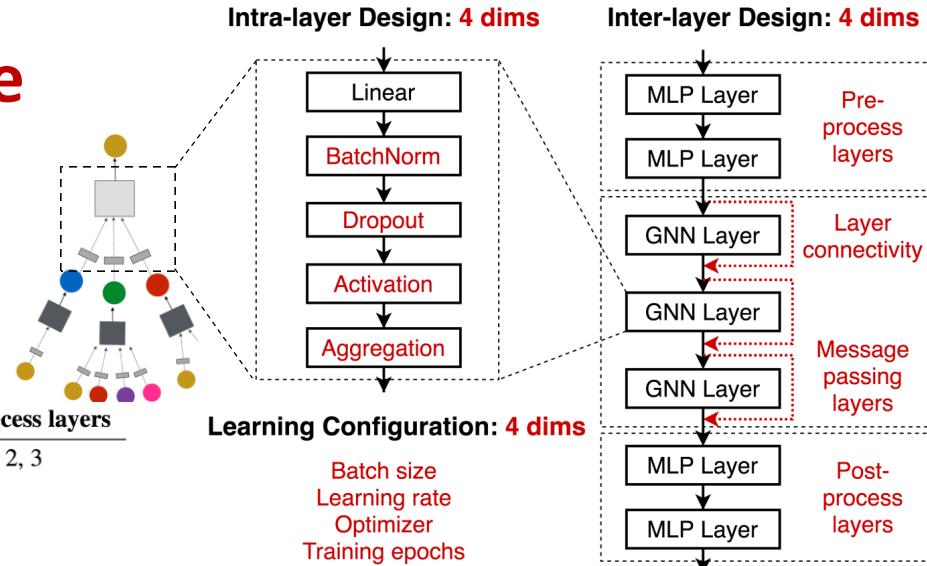
■ Learning configuration

Batch size	Learning rate	Optimizer	Training epochs
16, 32, 64	0.1, 0.01, 0.001	SGD, ADAM	100, 200, 400

■ In total: 315K possible designs

■ Our Purpose:

- We don't want to (and we cannot) cover all the possible designs
- **A mindset transition:** We want to demonstrate that **studying a design space is more effective than studying individual GNN designs**



A General GNN Task Space

- **Categorizing GNN tasks**
 - **Common practice:** node / edge / graph level task
 - Reasonable but not precise
 - **Node prediction:** predict **clustering coefficient** vs. predict a **node's subject area in a citation networks** – **completely different task**
 - But creating a precise taxonomy of GNN tasks is very hard!
 - **Subjective;** **Novel GNN tasks** can always emerge
- **Our innovation: a quantitative task similarity metric**
 - **Purpose:** **understand GNN tasks, transfer the best GNN models across tasks**

A General GNN Task Space

- Quantitative task similarity metric
 - 1) Select “anchor” models (M_1, \dots, M_5)
 - 2) Characterize a task by ranking the performance of anchor models
 - 3) Tasks with similar rankings are considered as similar

Task Similarity Metric

	Anchor Model Performance ranking					Similarity to Task A
Task A	M_1	M_2	M_3	M_4	M_5	1.0
Task B	M_1	M_3	M_2	M_4	M_5	0.8
Task C	M_5	M_1	M_4	M_3	M_2	-0.4

Task A is similar to Task B

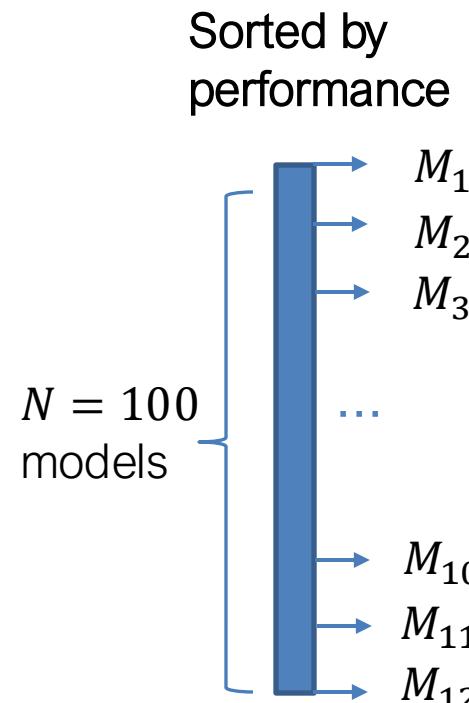
Task A is not similar to Task C

- How do we select the anchor models?

A General GNN Task Space

■ Selecting the anchor models

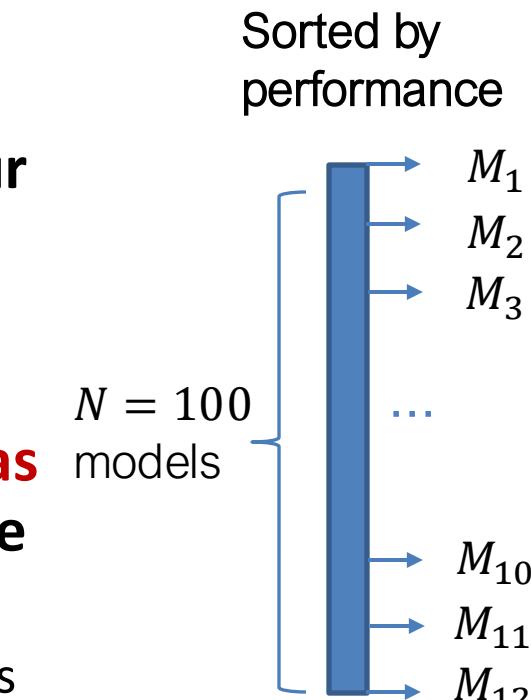
- 1) Select a small dataset
 - E.g., node classification on Cora
- 2) Randomly **sample N models from our design space**, run on the dataset
 - E.g., we sample 100 models
- 3) Sort these models based on their performance:



A General GNN Task Space

■ Selecting the anchor models

- 1) Select a small dataset
 - E.g., node classification on Cora
- 2) Randomly **sample N models from our design space**, run on the dataset
 - E.g., we sample 100 models
- 3) Sort these models based on their performance: **evenly select M models as the anchor models**, whose performance range from the worst to the best
 - E.g., we sample 12 models in our experiments
- **Goal: Cover a wide spectrum of models:**
A bad model in one task could be great for another task



A General GNN Task Space

■ We collect 32 tasks: node / graph classification

Task name
node AMAZON COMPUTERS N/A N/A
node AMAZON PHOTO N/A N/A
node CITESEER N/A N/A
node COAUTHORCS N/A N/A
node COAUTHORPHYSICS N/A N/A
node CORA N/A N/A
node scalefree-clustering-pagerank
node scalefree-const-clustering
node scalefree-const-pagerank
node scalefree-onehot-clustering
node scalefree-onehot-pagerank
node scalefree-pagerank-clustering
node smallworld-clustering-pagerank
node smallworld-const-clustering
node smallworld-const-pagerank
node smallworld-onehot-clustering
node smallworld-onehot-pagerank
node smallworld-pagerank-clustering
graph PROTEINS N/A N/A
graph BZR N/A N/A
graph COX2 N/A N/A
graph DD N/A N/A
graph ENZYMES N/A N/A
graph IMDB N/A N/A
graph scalefree-clustering-path
graph scalefree-const-path
graph scalefree-onehot-path
graph scalefree-pagerank-path
graph smallworld-clustering-path
graph smallworld-const-path
graph smallworld-onehot-path
graph smallworld-pagerank-path
graph ogbg-molhiv N/A N/A

(We include link prediction results in the Appendix)

6 Real-world node classification tasks

12 Synthetic node classification tasks

Predict node properties:

- Clustering coefficient
- PageRank

6 Real-world graph classification tasks

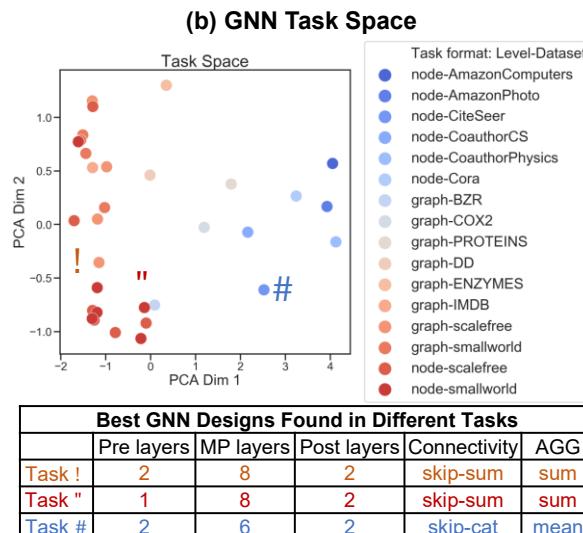
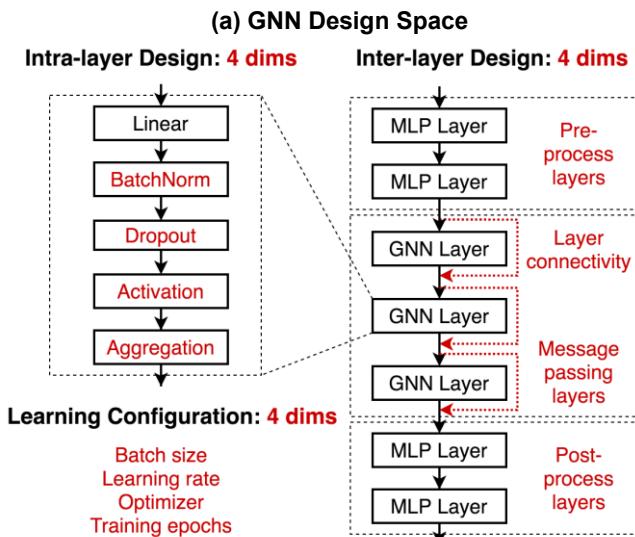
8 Synthetic graph classification tasks

Predict graph properties:

- Average path length

GNN Design Space

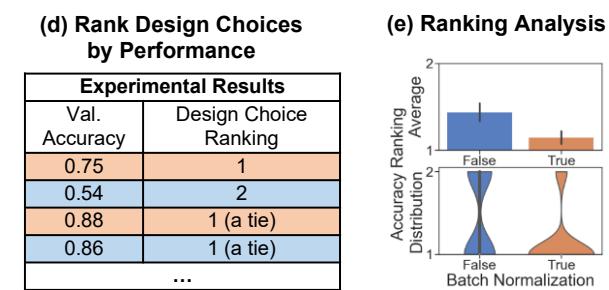
- Systematic investigation of:
 - General guidelines for GNN design
 - Understandings of GNN tasks
 - Transferring best GNN designs across tasks



(c) Controlled Random Search

GNN Design Space					GNN Task Space	
BatchNorm	Act	...	MP layers	Connectivity	level	dataset
True	relu	...	8	skip_sum	node	CiteSeer
False	relu	...	8	skip_sum	node	CiteSeer
True	relu	...	2	skip_cat	graph	BZR
False	relu	...	2	skip_cat	graph	BZR

... (truncated)



Evaluating GNN Designs

- **Evaluating a design dimension:**
 - “Is BatchNorm generally useful for GNNs?”
- **The common practice:**
 - (1) Pick one model (e.g., a 5-layer 64-dim GCN)
 - (2) Compare two models, with BN = True / False
- **Our approach:**
 - Note that **we have defined** 315K (models) * 32 (tasks) \approx **10M model-task combinations**
 - **(1) Sample from 10M possible model-task combinations**
 - **(2) Rank the models** with BN = True / False
- How do we make it **scalable & convincing?**

Evaluating GNN Designs

- Evaluating a design dimension: Controlled random search
 - a) Sample random model-task configurations, perturb BatchNorm = [True, False]
 - Here we control the computational budget for all the models

(a) Controlled Random Search

GNN Design Space					GNN Task Space	
BatchNorm	Activation	...	Message layers	Layer Connectivity	Task level	dataset
True	relu	...	8	skip_sum	node	CiteSeer
False	relu	...	8	skip_sum	node	CiteSeer
True	relu	...	2	skip_cat	graph	BZR
False	relu	...	2	skip_cat	graph	BZR
...						
True	prelu	...	4	stack	graph	scale free
False	prelu	...	4	stack	graph	scale free

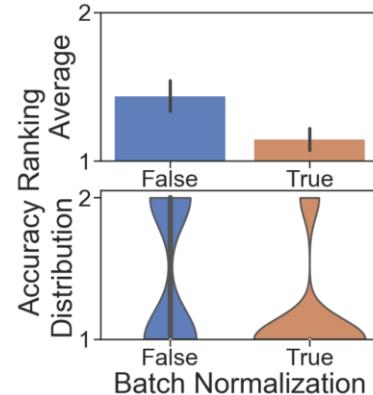
Evaluating GNN Designs

- **b) Rank** BatchNorm = [True, False] by their performance (lower ranking is better)
- **c) Plot Average / Distribution of the ranking** of BatchNorm = [True, False]

(b) Rank Design Choices by Performance

GNN Design Space		Experimental Results	
BatchNorm		Val. Accuracy	Design Choice Ranking
True		0.75	1
False		0.54	2
True		0.88	1 (a tie)
False		0.88	1 (a tie)
True		0.89	1
False		0.36	2

(c) Ranking Analysis

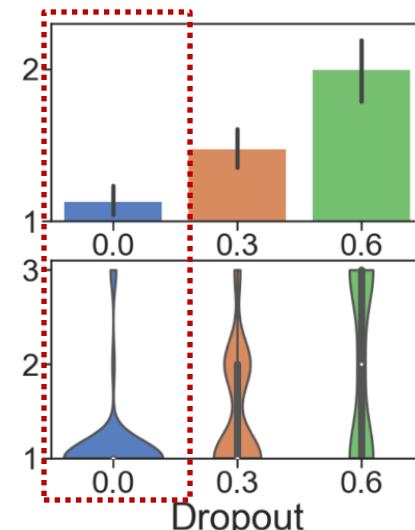
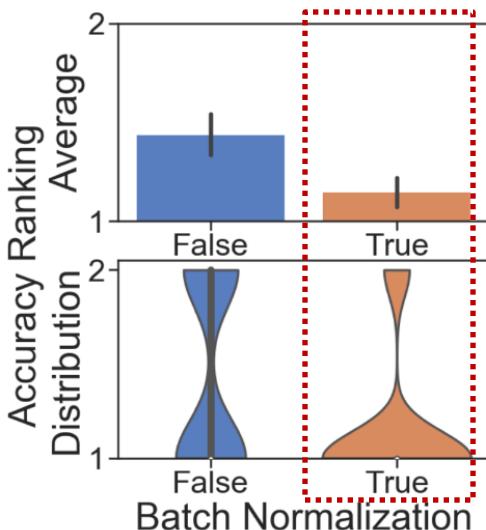


- **Summary:** Convincingly evaluate any new design dimension, e.g., evaluate a new GNN layer we propose

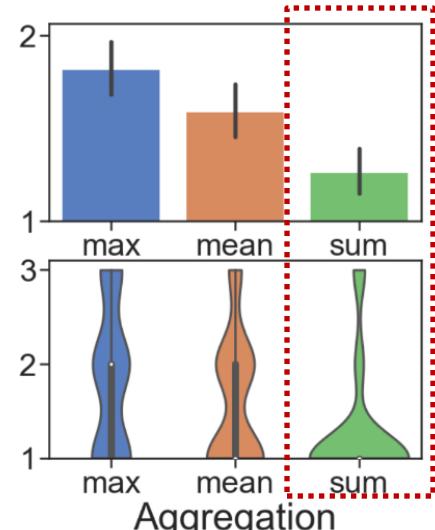
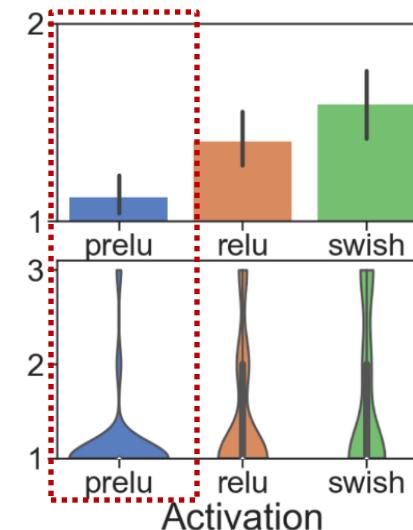
Results 1: A Guideline for GNN Design

- Certain design choices exhibit **clear advantages**
 - Intra-layer designs:

Explanation:
GNNs are hard to optimize



Explanation:
This is our new finding!



Explanation:
GNNs experience
underfitting more often

Explanation:
Sum is the most expressive
aggregator

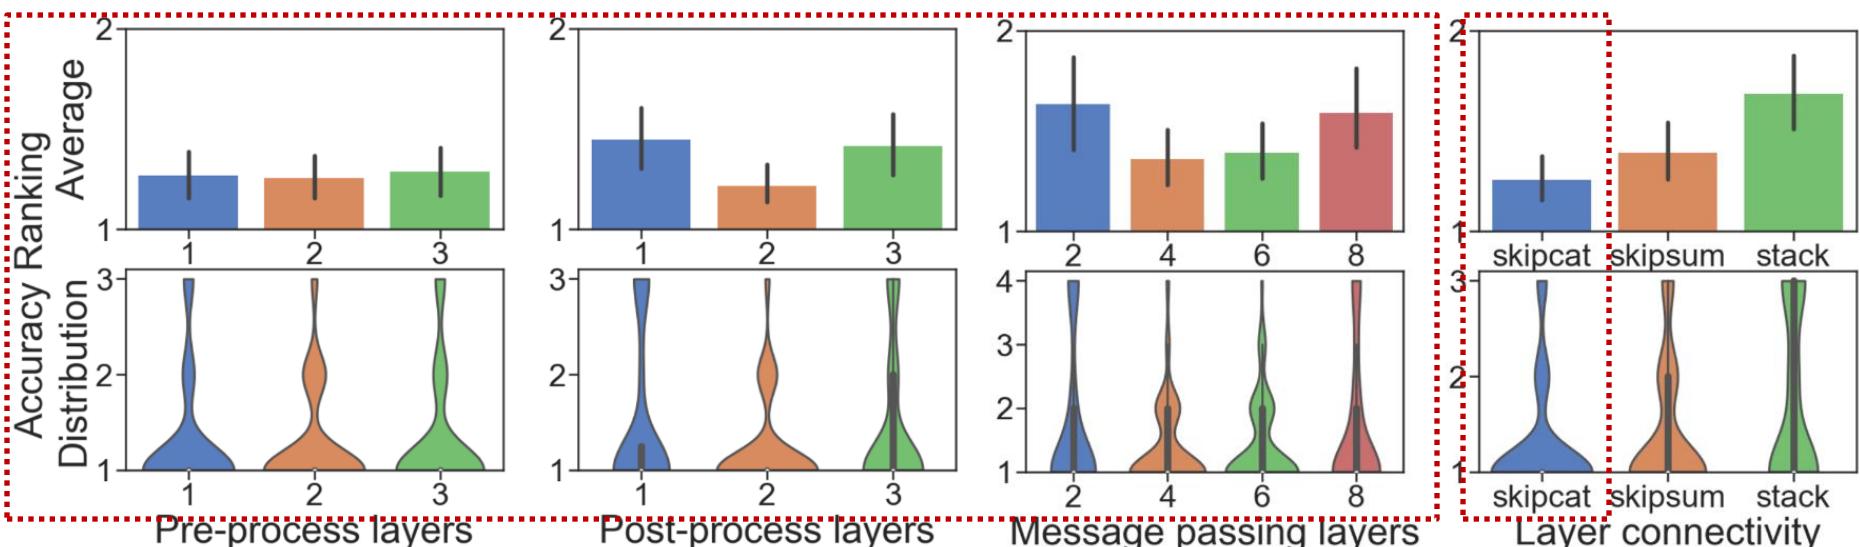
Results 1: A Guideline for GNN Design

- Certain design choices exhibit **clear advantages**

- **Inter-layer designs**

Optimal number of layers is hard to decide

Highly dependent on the task



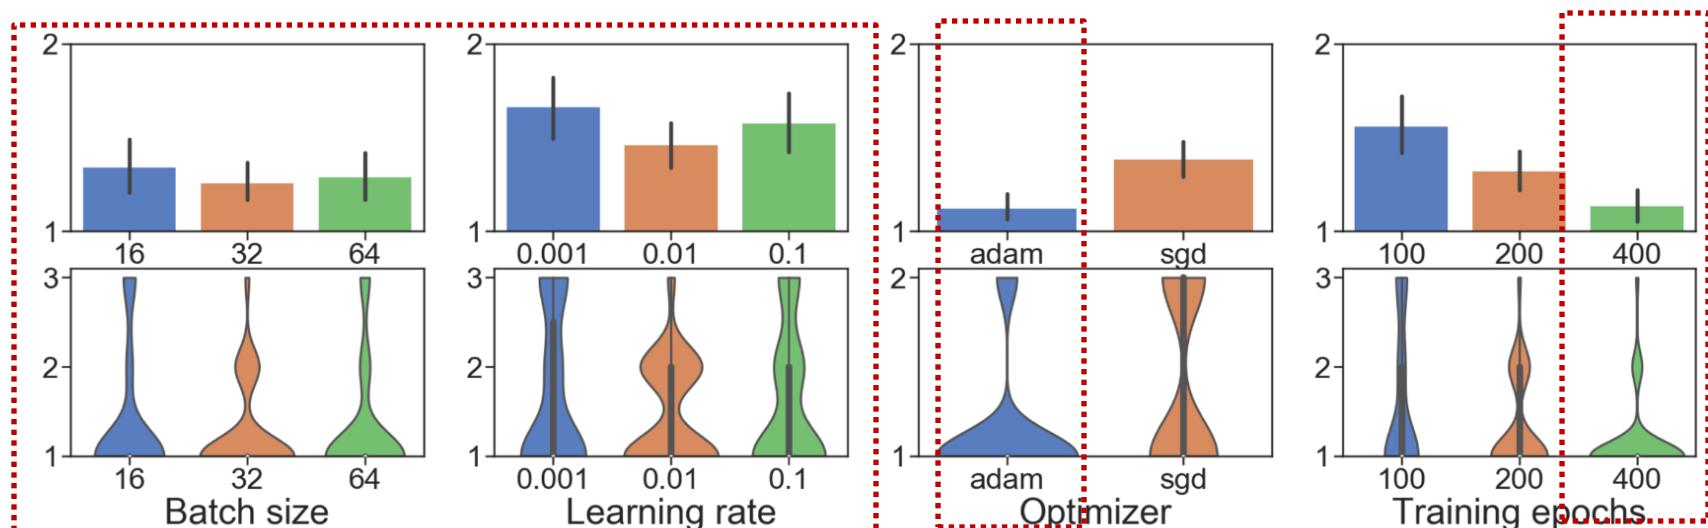
Explanation:
Skip connection enable hierarchical node representation

Results 1: A Guideline for GNN Design

- Certain design choices exhibit **clear advantages**
 - Learning configurations

Optimal batch size and learning rate is hard to decide

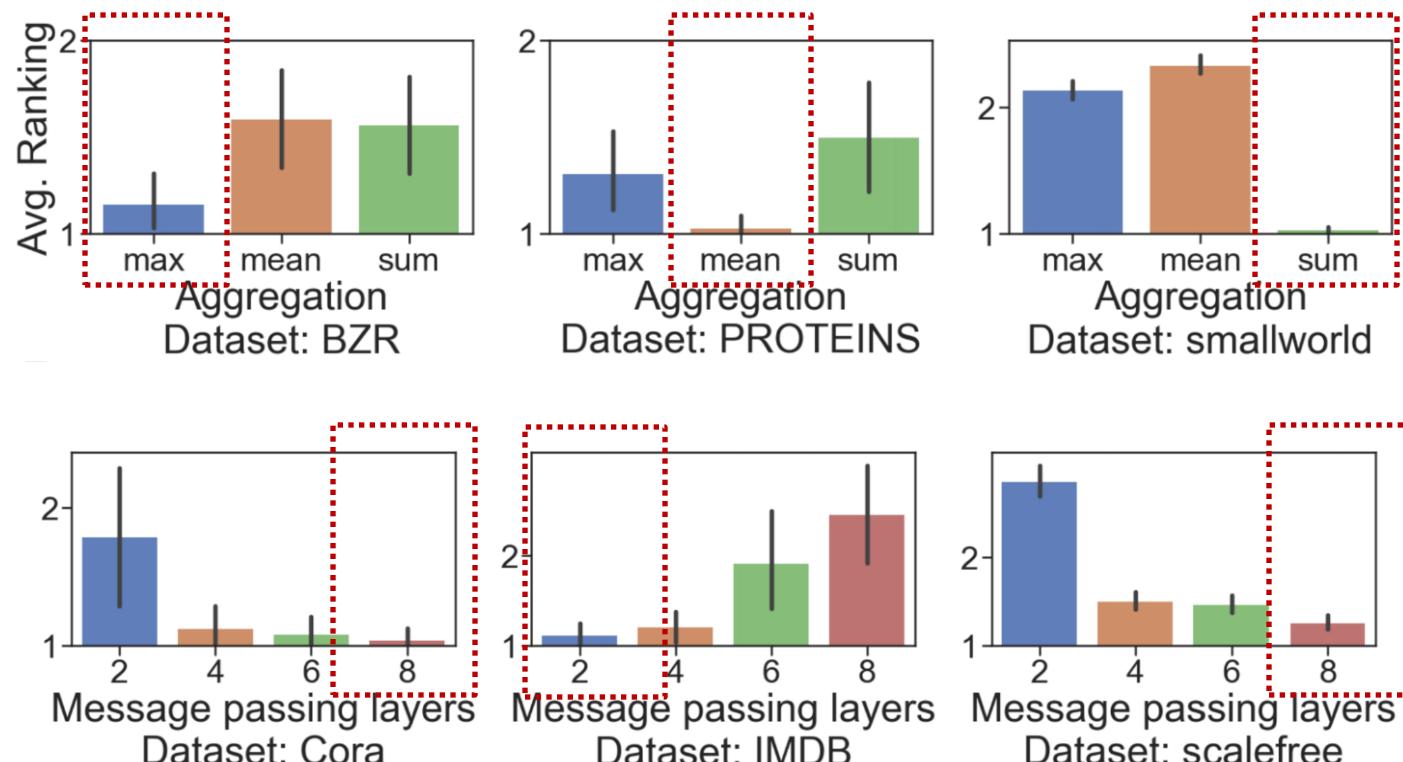
Highly dependent on the task



Explanation:
Adam is more robust
More training epochs is better

Results 2: Understanding GNN Tasks

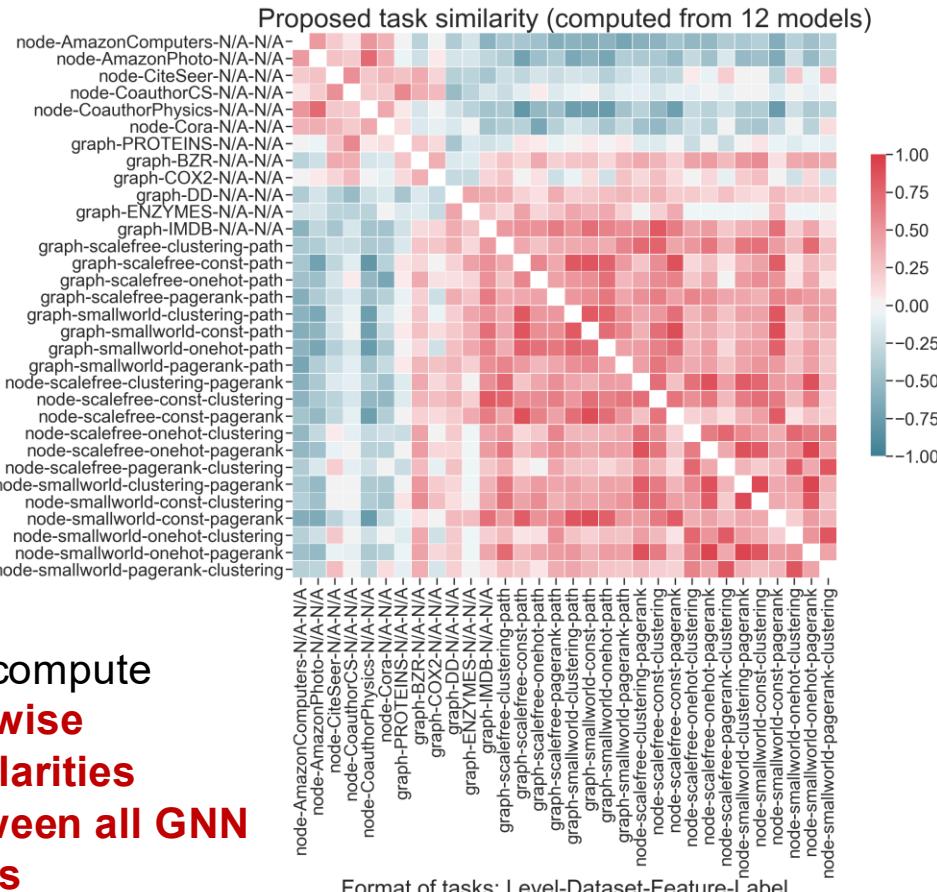
- Best GNN designs in different tasks **vary significantly**
 - Motivate that studying a task space is crucial



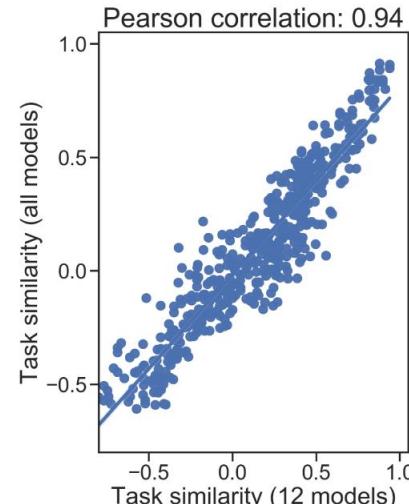
Results 2: Understanding GNN Tasks

■ Build a GNN task space

Recall how we **compute task similarity**



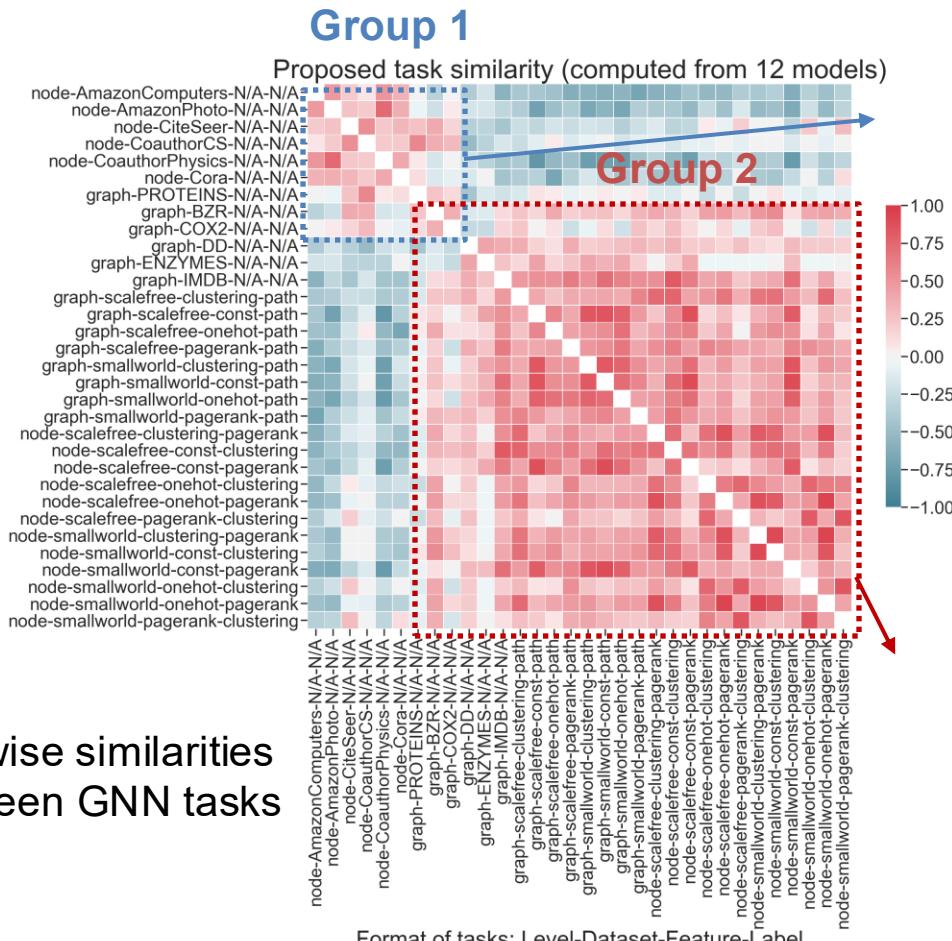
	Anchor Model Performance ranking					Similarity to Task A	
	Task A	M_1	M_2	M_3	M_4	M_5	
Task A		1.0					1.0
Task B		M_1	M_3	M_2	M_4	M_5	0.8
Task C		M_5	M_1	M_4	M_3	M_2	-0.4



Task similarity computation is cheap:
Using **12 anchor models** is a good approximation!

Results 2: Understanding GNN Tasks

GNN task space is informative



Group 1:

Tasks rely on **feature information**
Node/graph classification tasks,
where **input graphs have high-dimensional features**

- Cora graph has 1000+ dim node feature

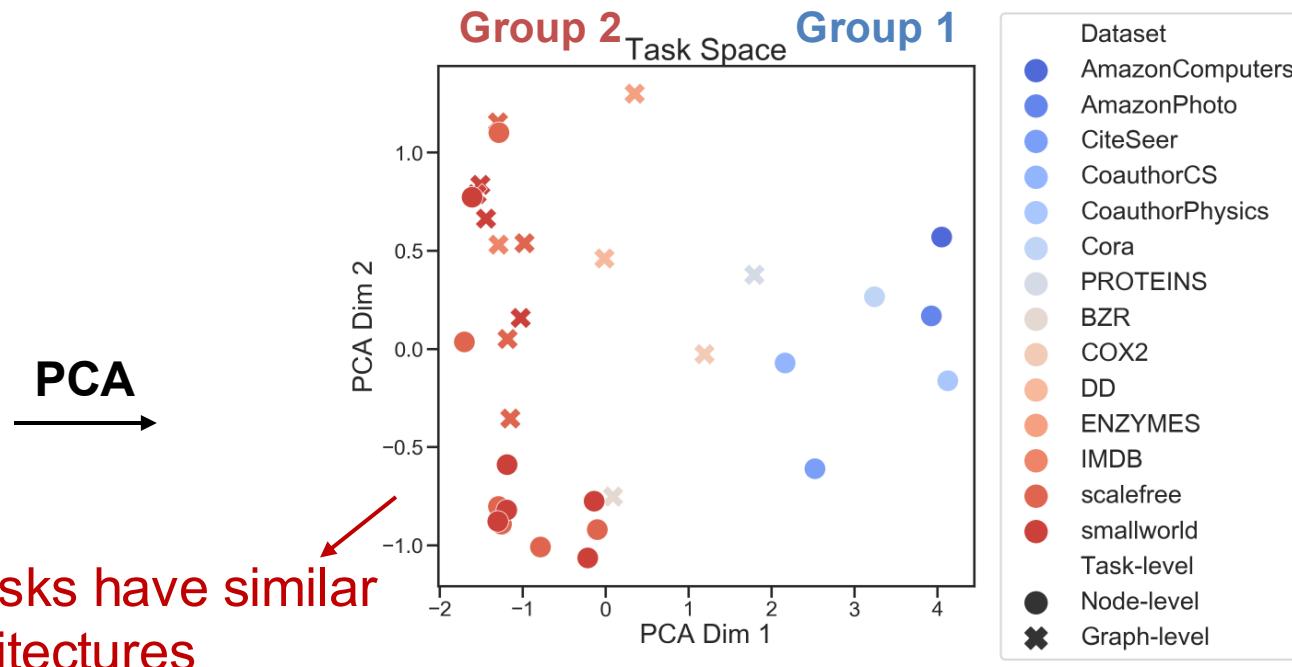
Group 2:

Tasks rely on **structural information**
Nodes have few features
Predictions are highly **dependent on graph structure**

- Predicting clustering coefficients

Results 2: Understanding GNN Tasks

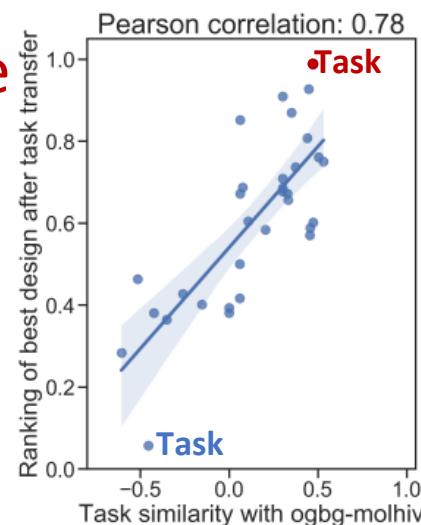
GNN task space is informative



	Best GNN Designs Found in Different Tasks				
	Pre layers	MP layers	Post layers	Connectivity	AGG
Task A	2	8	2	skip-sum	sum
Task B	1	8	2	skip-sum	sum
Task C	2	6	2	skip-cat	mean

Results 3: Transfer to Novel Tasks

- **Case study:** generalize best models to **unseen** OGB ogbg-molhiv task:
 - **ogbg-molhiv is unique from other tasks:** 20x larger, imbalanced (1.4% positive) and requires out-of-distribution generalization
- **Concrete steps for applying to a novel task:**
 - **Step 1:** Measure 12 anchor model performance on the new task
 - **Step 2:** Compute similarity between the new task and existing tasks
 - **Step 3:** Recommend the best designs from existing tasks with high similarity



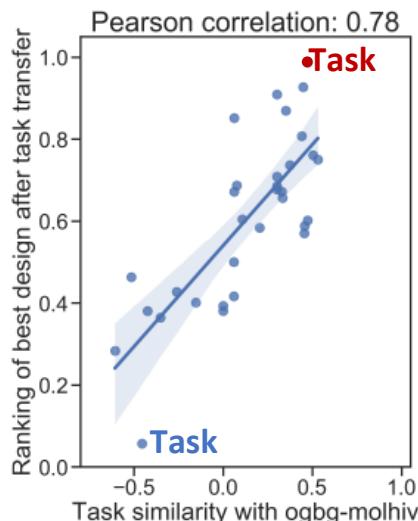
Results 3: Transfer to Novel Tasks

- Our task space can guide best model transfer to novel tasks!

We pick 2 tasks:

Task A: Similar to OGB

Task B: Not similar to OGB



Findings:

Transfer the best model from Task A achieves SOTA on OGB

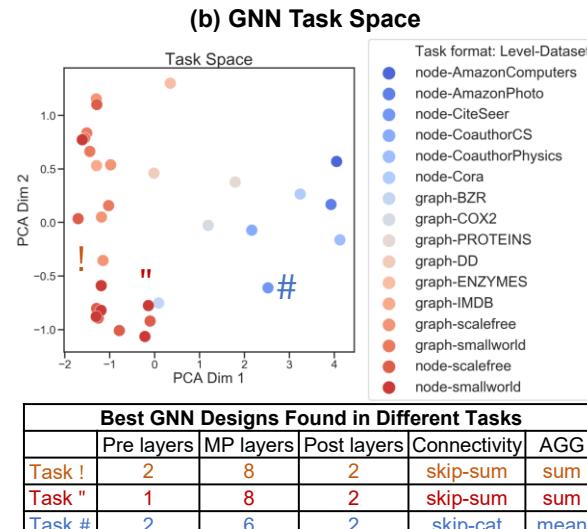
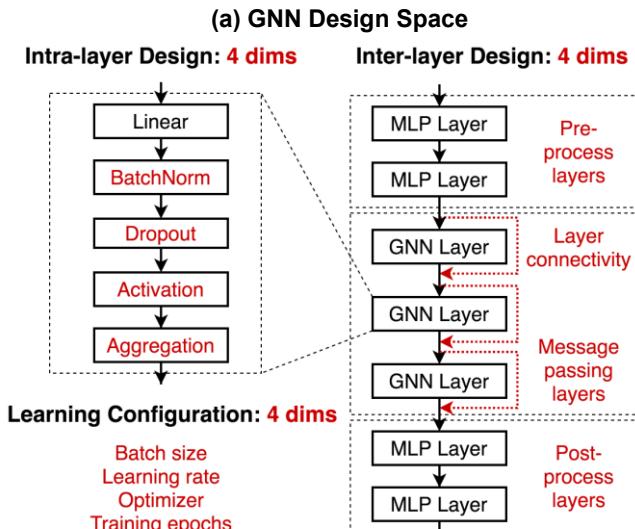
Transfer the best model from Task B performs badly on OGB

	Task A: graph-scalefree-const-path	Task B: node-CoauthorPhysics
Best design in our design space	(1, 8, 3, skipcat, sum)	(1, 4, 2, skipcat, max)
Task Similarity with ogbg-molhiv	0.47	-0.61
Performance after transfer to ogbg-molhiv	0.785	0.736

Previous SOTA: 0.771

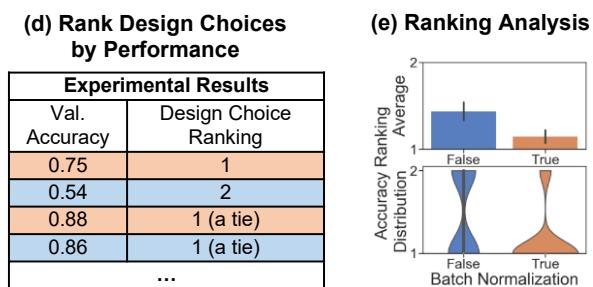
GNN Design Space: Summary

- Systematic investigation of:
 - General guidelines for GNN design
 - Understandings of GNN tasks
 - Transferring best GNN designs across tasks
 - **GraphGym:** Easy-to-use **code platform for GNN**



(c) Controlled Random Search

GNN Design Space				GNN Task Space		
BatchNorm	Act	... MP layers	Connectivity	level	dataset	
True	relu	...	8	skip_sum	node	CiteSeer
False	relu	...	8	skip_sum	node	CiteSeer
True	relu	...	2	skip_cat	graph	BZR
False	relu	...	2	skip_cat	graph	BZR
...						

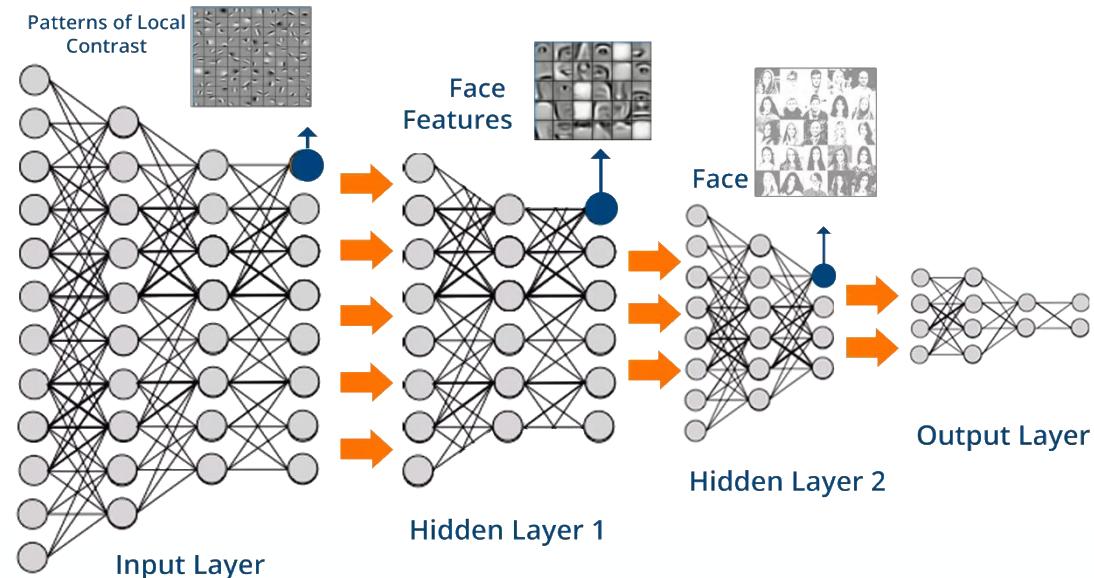
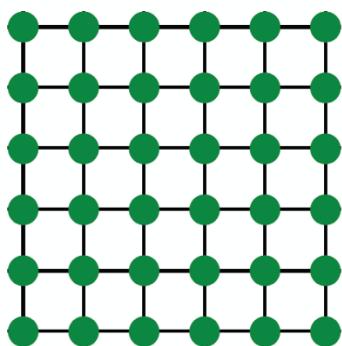


CS224W: Wrap-Up

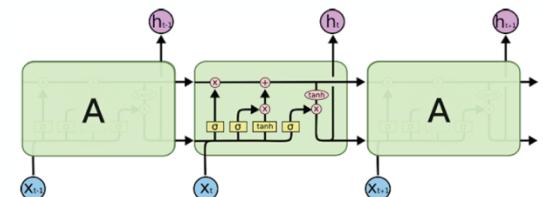
CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>



Modern ML Toolbox



Text/Speech



Modern deep learning toolbox is designed
for simple sequences & grids

This Course

How can we develop neural networks
that are much more broadly
applicable?

Graphs are the new frontier
of deep learning

Graphs and Relational Data

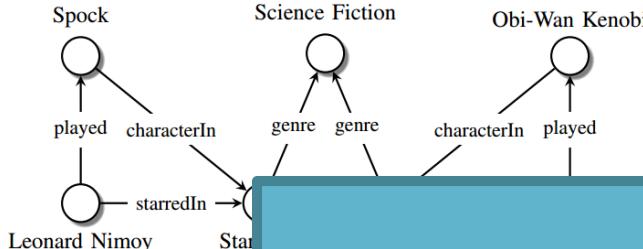


Image credit:

[Know](#)

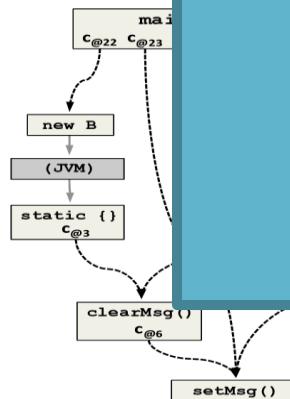


Image credit: ResearchGate

Code Graphs

Main question:

How do we take advantage of relational structure for better prediction?

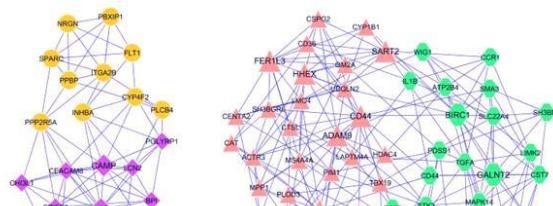


Image credit: MDPI

Molecules

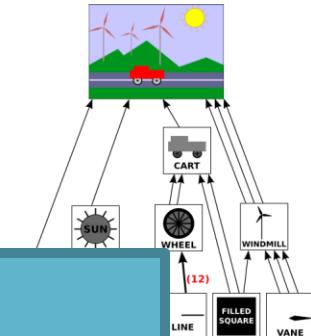


Image credit: [math.hws.edu](#)

Graphs

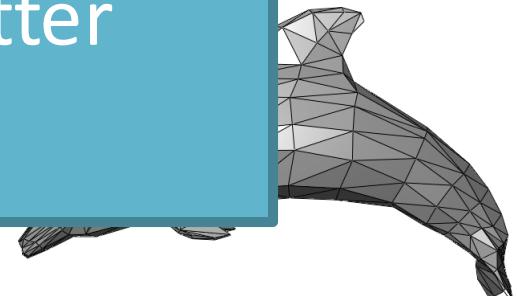
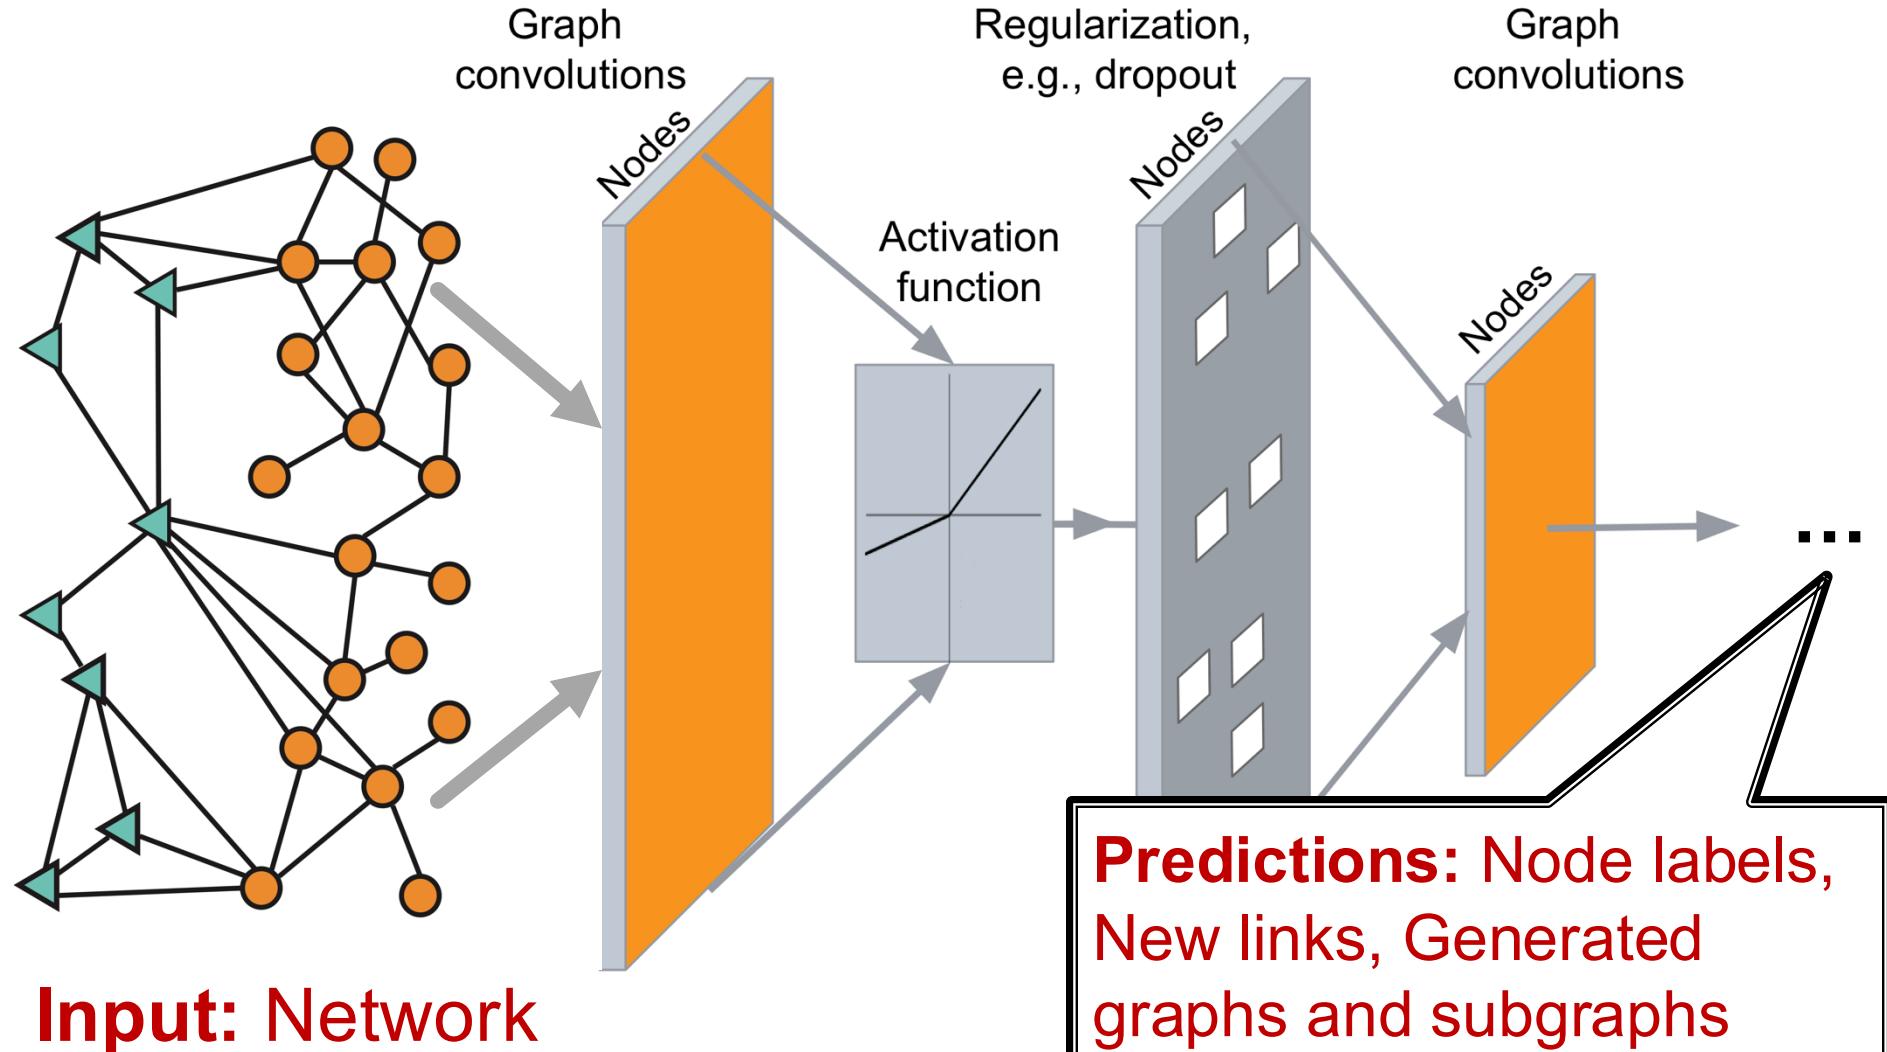


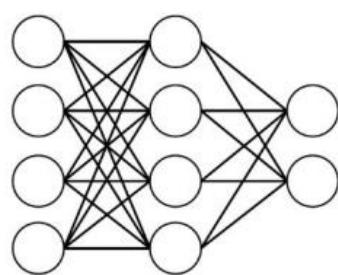
Image credit: Wikipedia

3D Shapes

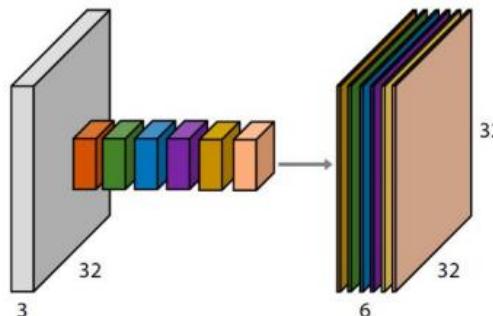
CS224W: Deep Learning in Graphs



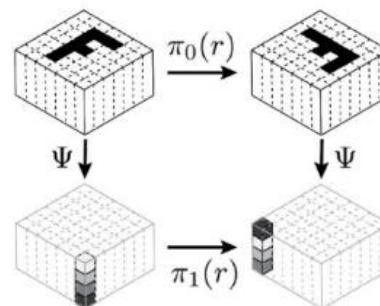
Models of Interest: Invariances



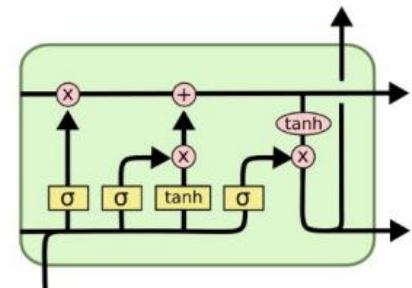
Perceptrons
Function regularity



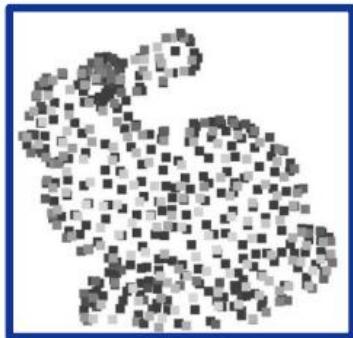
CNNs
Translation



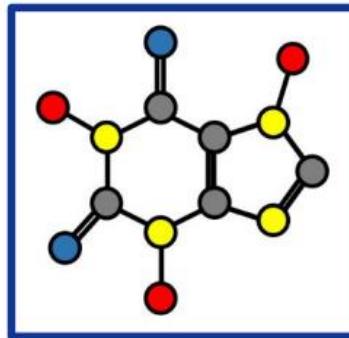
Group-CNNs
Translation+Rotation,
Global groups



LSTMs
Time warping



DeepSets / Transformers
Permutation



GNNs
Permutation



Intrinsic CNNs
Isometry / Gauge choice

The Bottom Line

- **There is exciting relational structure in many many real-world problems**
 - Molecules/Proteins as strings vs. graphs
 - Relational databases (primary-foreign key structure)
- **Identifying and harnessing this relational structure leads to better predictions**
 - AlphaFold
 - Biomedicine
 - Recommender systems

You learned a lot!

- **Theory:**
 - Models, architectures, approaches
- **Practice:**
 - Collab notebooks
 - Homeworks
- **Creative research:**
 - Course project
- **The real-world use cases and applications**

What Next?

- **Project write-ups:**

- Thurs Dec. 12, Midnight **(11:59PM)** Pacific Time

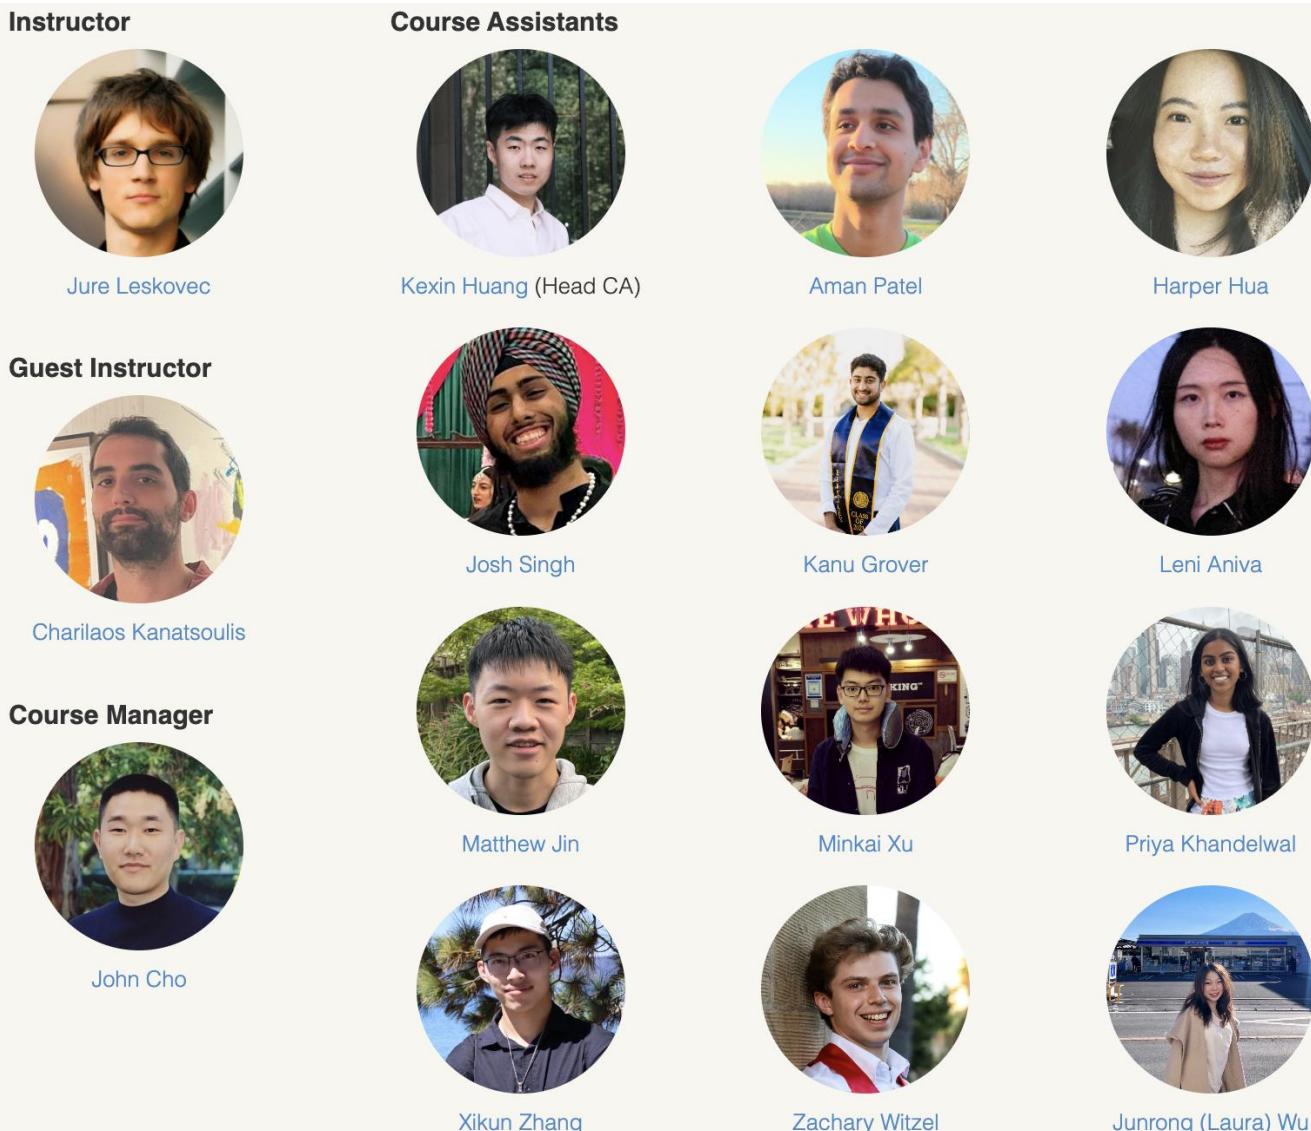
No late days!

- **Courses:**

- **CS246: Mining Massive Datasets (Winter)**

- Data Mining & Machine Learning for big data
 - (big==doesn't fit in memory/single machine)
 - Fast clever algorithms for real-world problems
 - Distributed data processing frameworks:
MapReduce, Spark

Thank you, team!!!



I am very proud of everyone!

- You Have Done a Lot!!!
- And (hopefully) learned a lot!!!
 - Answered questions and proved many interesting results
 - Implemented a number of methods
 - And are doing excellently on the project!

Thank You for the
Hard Work!!!