

Note to other teachers and users of these slides: We would be delighted if you found our material useful for giving your own lectures. Feel free to use these slides verbatim, or to modify them to fit your own needs. If you make use of a significant portion of these slides in your own lecture, please include this message, or a link to our web site: <http://cs224w.Stanford.edu>

Stanford CS224W: How Expressive are Graph Neural Networks?


CS224W: Machine Learning with Graphs
Charilaos Kanatsoulis and Jure Leskovec, Stanford
University
<http://cs224w.stanford.edu>



Announcements

- **Colab 1** due today at 11:59PM
- **Homework 1** due in 1 week (Thursday 10/16)
- **Colab 2** will be released today at 3PM
 - Due on Thursday 10/23 (2 weeks from now)
- **Project Proposal** is due Tuesday 10/21
 - Details on course website
 - To do a solo project, make a private post on Ed with your idea to get approval from course staff

Recap: GNN Training Pipeline



Implementation resources:

[PyG](#) provides core modules for this pipeline


[GraphGym](#) further implements the full pipeline to facilitate GNN design

Information modalities in Graphs



Image credit: [Medium](#)

Social Networks



Citation Networks




Image credit: [Science](#)

Economic Networks





Image credit: [MDPI](#)

Molecules



Why is Graph Deep Learning Hard?

Graphs are complex.


- Complex topological structure (*i.e.*, no spatial locality like grids)



Networks



Images



Text

- No fixed node ordering or reference point



Theory of GNNs

How powerful are GNNs?

- Many GNN models have been proposed (e.g., GCN, GAT, GraphSAGE, design space).
- What is the expressive power (ability to distinguish different graph structures) of these GNN models?
- How to design a maximally expressive GNN model?

Graph Isomorphism problem


- A classical expressivity test is graph isomorphism.
 - Given a pair of graphs, can we tell if they are isomorphic?



- Open problem (a polynomial time algorithm does not exist, has not been proven NP-complete).


A simpler problem

- Given a pair of nodes with different neighborhood structures. Is there a GNN that can always tell them apart?





Background: A Single GNN Layer

- We focus on message passing GNNs:
 - (1) Message: each node computes a message
$$\mathbf{m}_u^{(l)} = \text{MSG}^{(l)}\left(\mathbf{h}_u^{(l-1)}\right), u \in \{N(v) \cup v\}$$
 - (2) Aggregation: aggregate messages from neighbors
$$\mathbf{h}_v^{(l)} = \text{AGG}^{(l)}\left(\left\{\mathbf{m}_u^{(l)}, u \in N(v)\right\}, \mathbf{m}_v^{(l)}\right)$$



Background: Many GNN Models



- Many GNN models have been proposed:
 - GCN, GraphSAGE, GAT, Design Space etc.



Different GNN models use different neural networks in the box



GNN Model Example (1)

- GCN (mean-pool) [Kipf and Welling ICLR 2017]



GNN Model Example (2)


- GraphSAGE (max-pool) [Hamilton et al. NeurIPS 2017]



MLP + element-wise max-pooling

Note: Node Colors

- We use node same/different **colors** to represent nodes with same/different features.
 - For example, the graph below assumes all the nodes share the same feature.




- **Key question:** How well can a GNN distinguish different graph structures?

Local Neighborhood Structures

- We specifically consider **local neighborhood structures** around each node in a graph.


- **Example:** Nodes 1 and 5

have **different** neighborhood structures because they have different node degrees.



Local Neighborhood Structures


- We specifically consider **local neighborhood structures** around each node in a graph.
 - **Example:** Nodes 1 and 4 both have the same node degree of 2. However, they still have **different** neighborhood structures because **their neighbors have different node degrees.**



Node 1 has neighbors of degrees 2 and 3.
Node 4 has neighbors of degrees 1 and 3.

Local Neighborhood Structures

- We specifically consider **local neighborhood structures** around each node in a graph.
 - **Example:** Nodes 1 and 2 have the **same** neighborhood structure because **they are symmetric within the graph.**



Node 1 has neighbors of degrees 2 and 3.

Node 2 has neighbors of degrees 2 and 3.



And even if we go a step deeper to 2nd hop neighbors, both nodes have the same degrees (Node 4 of degree 2)

Local Neighborhood Structures

- **Key question:** Can GNN node embeddings distinguish different node's local neighborhood structures?
 - If so, when? If not, when will a GNN fail?
- **Next:** We need to understand how a GNN captures local neighborhood structures.
 - Key concept: Computational graph


Computational Graph (1)

- In each layer, a GNN aggregates neighboring node embeddings.
- A GNN generates node embeddings through a computational graph defined by the neighborhood.
 - Ex: Node 1's computational graph (2-layer GNN)




Computational Graph (2)

- Ex: Nodes 1 and 2's computational graphs.



Computational Graph (3)


- Ex: Nodes 1 and 2's computational graphs.
- But GNN only sees node features (not IDs):



Computational Graph (4)

- A GNN will generate the same embedding for nodes 1 and 2 because:
 - Computational graphs are the same.
 - Node features (colors) are identical.


Note: GNN does not care about node ids, it just aggregates features vectors of different nodes.



GNN won't be able to distinguish nodes 1 and 2


Computational Graph

- In general, different local neighborhoods define different computational graphs



Computational Graph


- Computational graphs are identical to **rooted subtree structures** around each node.



Computational Graph


- GNN's node embeddings capture **rooted subtree structures**.
- Most expressive GNN maps different **rooted subtrees** into different node embeddings (represented by different colors).

Embedding




Recall: Injective Function

- **Function** $f: X \rightarrow Y$ is **injective** if it maps different elements into different outputs.
- **Intuition:** f retains all the information about input.




How Expressive is a GNN?

- Most expressive GNN should map subtrees to the node embeddings **injectively**.




How Expressive is a GNN?

- **Key observation:** Subtrees of the same depth can be recursively characterized from the leaf nodes to the root nodes.




How Expressive is a GNN?

- If each step of GNN's aggregation **can fully retain the neighboring information**, the generated node embeddings can distinguish different rooted subtrees.



How Expressive is a GNN?

- In other words, most expressive GNN would use an **injective neighbor aggregation** function at each step.
 - Maps different neighbors to different embeddings.




How Expressive is a GNN?


■ Summary so far

- To generate a node embedding, GNNs use a computational graph corresponding to a **subtree rooted around each node**.

Input graph



Computational graph
= Rooted subtree



Using injective neighbor aggregation → distinguish different subtrees

- GNN can fully distinguish different subtree structures if **every step of its neighbor aggregation is injective**.

Stanford CS224W: Designing the Most Powerful Graph Neural Network

CS224W: Machine Learning with Graphs

Jure Leskovec, Stanford University

<http://cs224w.stanford.edu>




Expressive Power of GNNs

- **Key observation:** Expressive power of GNNs can be characterized by that of neighbor aggregation functions they use.
 - A more expressive aggregation function leads to a more expressive GNN.
 - Injective aggregation function leads to the most expressive GNN.
- **Next:**
 - Theoretically analyze expressive power of aggregation functions.


Neighbor Aggregation

- **Observation:** Neighbor aggregation can be abstracted as **a function over a multi-set** (a set with repeating elements).




Equivalent

A double-headed grey arrow indicating equivalence between the two representations. To the left of the arrow is the 'Neighbor aggregation' diagram. To the right is a diagram showing a red bracket-like shape enclosing two yellow circles, with a red line extending upwards from the top of the bracket to a single red peak, representing a multi-set function.



Multi-set function

Examples of multi-set



Same color indicates the same features.

Neighbor Aggregation


- **Next:** We analyze aggregation functions of two popular GNN models
 - **GCN** (mean-pool) [Kipf & Welling, ICLR 2017]
 - Uses **element-wise** mean pooling over neighboring node features
$$\text{Mean}(\{x_u\}_{u \in N(v)})$$
 - **GraphSAGE** (max-pool) [Hamilton et al. NeurIPS 2017]
 - Uses **element-wise** max pooling over neighboring node features
$$\text{Max}(\{x_u\}_{u \in N(v)})$$

Neighbor Aggregation: Case Study

■ GCN (mean-pool) [Kipf & Welling ICLR 2017]

- Take **element-wise mean**, followed by linear function and ReLU activation, i.e., $\max(0, x)$.
- **Theorem** [Xu et al. ICLR 2019]
 - GCN's aggregation function **cannot distinguish different multi-sets with the same color proportion**.

Failure case



■ Why?

Neighbor Aggregation


- For simplicity, we assume node features (colors) are represented by **one-hot encoding**.
 - Example: If there are two distinct colors:

$$\text{Yellow Circle} = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \quad \text{Blue Circle} = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

- This assumption is sufficient to illustrate how GCN fails.

Neighbor Aggregation: Case Study


- **GCN (mean-pool)** [Kipf & Welling ICLR 2017]
 - Failure case illustration



Neighbor Aggregation: Case Study

- **GraphSAGE (max-pool)** [Hamilton et al. NeurIPS 2017]
 - Apply an MLP, then take **element-wise max**.
 - **Theorem** [Xu et al. ICLR 2019]
 - GraphSAGE's aggregation function cannot distinguish different multi-sets with the same set of distinct colors.

Failure case



- **Why?**

Neighbor Aggregation: Case Study


- **GraphSAGE (max-pool)** [Hamilton et al. NeurIPS 2017]
 - Failure case illustration

The same outputs!

Element-wise-
max-pool

$$\begin{pmatrix} 1 \\ 1 \end{pmatrix} \longleftrightarrow \begin{pmatrix} 1 \\ 1 \end{pmatrix} \longleftrightarrow \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$

For simplicity,
assume the one-
hot encoding
after **MLP**.



Summary So Far

- We analyzed the **expressive power of GNNs**.
- **Main takeaways:**
 - Expressive power of GNNs can be characterized by that of the neighbor aggregation function.
 - Neighbor aggregation is a function over multi-sets (sets with repeating elements)
 - GCN and GraphSAGE's aggregation functions fail to distinguish some basic multi-sets; hence **not injective**.
 - Therefore, GCN and GraphSAGE are **not** maximally powerful GNNs.

Designing Most Expressive GNNs

- **Our goal:** Design maximally powerful GNNs in the class of message-passing GNNs.
- This can be achieved by designing **injective neighbor aggregation function over multisets**.
- Here, we design a **neural network** that can model **injective multiset function**.

Injective Multi-Set Function

Theorem [Xu et al. ICLR 2019]



Any injective multi-set function can be expressed as:

$$\text{Some non-linear function} \xrightarrow{\quad} \Phi \left(\sum_{x \in S} f(x) \right)$$

Some non-linear function

Sum over multi-set

S : multi-set



$$\Phi \left(f[\text{Yellow Circle}] + f[\text{Blue Circle}] + f[\text{Blue Circle}] \right)$$

Injective Multi-Set Function

Proof Intuition: [Xu et al. ICLR 2019]

f produces one-hot encodings of colors. Summation of the one-hot encodings retains all the information about the input multi-set.

$$\Phi \left(\sum_{x \in S} f(x) \right)$$

Example: $\Phi \left[f([\bullet]) + f([\bullet]) + f([\bullet]) \right]$


One-hot $\begin{pmatrix} 1 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 \\ 2 \end{pmatrix}$

Universal Approximation Theorem

- How to model Φ and f in $\Phi(\sum_{x \in S} f(x))$?
- We use a Multi-Layer Perceptron (MLP).
- **Theorem: Universal Approximation Theorem**

[Hornik et al., 1989]

- 1-hidden-layer MLP with sufficiently-large hidden dimensionality and appropriate non-linearity $\sigma(\cdot)$ (including ReLU and sigmoid) can **approximate any continuous function to an arbitrary accuracy**.



Injective Multi-Set Function

- We have arrived at a **neural network** that can model any injective multiset function.

$$\text{MLP}_\Phi \left(\sum_{x \in S} \text{MLP}_f(x) \right)$$

- In practice, MLP hidden dimensionality of 100 to 500 is sufficient.

Most Expressive GNN

- **Graph Isomorphism Network (GIN)** [Xu et al. ICLR 2019]




- Apply an MLP, element-wise sum, followed by another MLP.

$$\text{MLP}_\Phi \left(\sum_{x \in S} \text{MLP}_f(x) \right)$$

- **Theorem** [Xu et al. ICLR 2019]
 - GIN's neighbor aggregation function is injective.
- **No failure cases!**
- **GIN is THE most expressive GNN in the class of message-passing GNNs we have introduced!**





Discussion: The Power of Pooling

Failure cases for mean and max pooling:



Colors represent feature values

Ranking by discriminative power:





Full Model of GIN

- **So far:** We have described the neighbor aggregation part of GIN.
- We now describe the full model of GIN by relating it to **WL graph kernel** (traditional way of obtaining graph-level features).
 - We will see how GIN is a “neural network” version of the WL graph kernel.

Graph Isomorphism problem

- A classical expressivity test is **Graph isomorphism**.
 - Given a pair of graphs, can we tell if they are isomorphic?



- Open problem (a polynomial time algorithm does not exist, has not been proven NP-complete)

Relation to WL Graph Kernel



Color refinement algorithm in WL kernel.

- **Given:** A graph G with a set of nodes V .
 - Assign an initial color $c^{(0)}(\nu)$ to each node ν .
 - Iteratively refine node colors by
$$c^{(k+1)}(\nu) = \text{HASH}\left(c^{(k)}(\nu), \{c^{(k)}(u)\}_{u \in N(\nu)}\right),$$
where HASH maps different inputs to different colors.
- After K steps of color refinement, $c^{(K)}(\nu)$ summarizes the structure of K -hop neighborhood



Color Refinement (1)

Example of color refinement given two graphs

- Assign initial colors





- Aggregate neighboring colors





Color Refinement (2)

Example of color refinement given two graphs

- Aggregated colors:



- Injectively HASH the aggregated colors






HASH table: **Injective!**

1,1	-->	2
1,11	-->	3
1,111	-->	4
1,1111	-->	5

Color Refinement (3)

Example of color refinement given two graphs

- Process continues until a stable coloring is reached
- Two graphs are considered **isomorphic** if they have the same set of colors.



The Complete GIN Model

- GIN uses a **neural network** to model the injective HASH function.

$$c^{(k+1)}(v) = \text{HASH}\left(c^{(k)}(v), \{c^{(k)}(u)\}_{u \in N(v)}\right)$$

- Specifically, we will model the injective function over the tuple:

$$(c^{(k)}(v), \{c^{(k)}(u)\}_{u \in N(v)})$$

Root node
features

Neighboring
node colors

The Complete GIN Model

Theorem (Xu et al. ICLR 2019)

Any injective function over the tuple

Root node
feature

$(c^{(k)}(v), \{c^{(k)}(u)\}_{u \in N(v)})$

Neighboring
node features

can be modeled as

$$\text{MLP}_\Phi \left((1 + \epsilon) \cdot \text{MLP}_f(c^{(k)}(v))) + \sum_{u \in N(v)} \text{MLP}_f(c^{(k)}(u)) \right)$$

where ϵ is a learnable scalar.

The Complete GIN Model

- If input feature $c^{(0)}(v)$ is represented as one-hot, **direct summation is injective**.

Example: $\Phi \left[\begin{array}{c} \text{Yellow circle} \\ + \\ \text{Blue circle} \\ + \\ \text{Blue circle} \end{array} \right]$

$$\begin{pmatrix} 1 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 \\ 2 \end{pmatrix}$$

- We only need Φ to ensure the injectivity.

$$\text{GINConv} \left(c^{(k)}(v), \left\{ c^{(k)}(u) \right\}_{u \in N(v)} \right) = \text{MLP}_\Phi \left((1 + \epsilon) \cdot c^{(k)}(v) + \sum_{u \in N(v)} c^{(k)}(u) \right)$$

Root node features Neighboring node features

This MLP can provide “one-hot” input feature for the next layer.

The Complete GIN Model

- GIN's node embedding updates
- Given: A graph G with a set of nodes V .
 - Assign an **initial vector** $c^{(0)}(\nu)$ to each node ν .
 - Iteratively update node vectors by

$$c^{(k+1)}(\nu) = \text{GINConv} \left(\left\{ c^{(k)}(\nu), \left\{ c^{(k)}(u) \right\}_{u \in N(\nu)} \right\} \right),$$

Differentiable color HASH function

where **GINConv** maps different inputs to different embeddings.

- After K steps of GIN iterations, $c^{(K)}(\nu)$ summarizes the structure of K -hop neighborhood.

GIN and WL Graph Kernel

- GIN can be understood as differentiable neural version of the WL graph Kernel:

	Update target	Update function
WL Graph Kernel	Node colors (one-hot)	HASH
GIN	Node embeddings (low-dim vectors)	GINConv



- Advantages of GIN over the WL graph kernel are:
 - Node embeddings are **low-dimensional**; hence, they can capture the fine-grained similarity of different nodes.
 - Parameters of the update function can be **learned for the downstream tasks**.

Expressive Power of GIN

- Because of the relation between GIN and the WL graph kernel, their expressive power is exactly the same.
 - If two graphs can be distinguished by GIN, they can be also distinguished by the WL kernel, and vice versa.
- How powerful is this?
 - WL kernel has been both theoretically and empirically shown to distinguish most of the real-world graphs [Cai et al. 1992].
 - Hence, GIN is also powerful enough to distinguish most of the real graphs!

Some bad news




- The WL kernel cannot distinguish between some basic graph structures, e.g.,



- The WL kernel cannot count important substructures, e.g., cycles and cliques.


Why do we care about cycles?

- The **bonds between different atoms** in chemical compounds form **cycles**.
- The **hexagon** is the **most common polygon** in organic compounds ⇒ yields stable materials.




Why cliques and quasi-cliques?

Fraudster Detection in e-Commerce Networks




[Hooi et al. 2016]

Detection of Protein Complexes in PPI Networks



[Vlaic et al. 2017]

Graph Compression




Dhulipala et al. 2017

Improving GNNs' Power

■ Can the expressive power of GNNs be improved?

- There are basic graph structures that existing GNN framework cannot distinguish, such as difference in cycles.



- GNNs' expressive power **can be improved** to resolve the above problem. [You et al. AAAI 2021, Li et al. NeurIPS 2020, Kanatsoulis et al. ICLR 2024]

Summary of the Lecture

- We design a neural network that can model an **injective multi-set function**.
- We use the neural network for neighbor aggregation function and arrive at **GIN---the most expressive GNN model**.
- The key is to use **element-wise sum pooling**, instead of mean-/max-pooling.
- GIN is closely related to the WL graph kernel.
- Both GIN and WL graph kernel can distinguish most of the real graphs!